

What Clarity Costs: Measuring Legible Robot Motion Against Path Cost and Constraint Satisfaction

Munawar Kazmi

Independent

United Kingdom

munawarsaeedkazmi1@gmail.com

Abstract

A robot moving towards one of several possible goals is unreadable for the first seconds of its motion, which is when a person nearby hesitates or steps into its path. Legible motion deviates early to resolve that ambiguity, and the deviation costs path length. What it can also cost is the constraint the robot was supposed to respect. We present a small judge-free benchmark that measures the three together: legibility in the form of Dragan et al., path cost against the exact optimum, and satisfaction of stated keep-out constraints. Eight worlds each carry machine-checked facts about their own geometry, optimal cost-to-go is computed exactly over a visibility graph, and no human rater or model judge appears anywhere in the scoring loop. Asked to produce legible motion under a stated path budget, two language models called all 80 of their trajectories legible; 25 of those were not physically possible. In the world built so that the clearest signal is the one the robot is not allowed to send, both models bought the clarity and entered the keep-out zone on every one of ten samples. Given four different path budgets, one model's median path cost moved by 0.012 and the other's did not move at all.

Keywords

legibility, human-robot interaction, motion planning, benchmarking, language models

1 Introduction

Legibility is not ours. Dragan et al. [2] formalised it: a legible trajectory is one that departs from the efficient path early enough for an observer to infer the goal sooner than efficiency alone would allow. What is missing from that literature is a price. Deviating for clarity moves the robot somewhere it would not otherwise go, and in a real room that somewhere is occupied.

The founding paper notices this. It reports that legibility will move a trajectory much closer to an obstacle in order to disambiguate between goals [3]. It does not measure it: obstacles enter as a soft term inside the objective, clearance is never reported, and constraint satisfaction is not an axis of any result. Legible navigation among moving agents does report proximity and time to collision beside legibility [1], but as properties of a policy rather than as a curve against a path budget.

This report contributes an instrument rather than a planner: a reproducible way to ask what a given trajectory pays, in path and in constraint satisfaction, for the clarity it buys. We use it to ask whether language models produce legible motion when told to.

2 The instrument

Everything is 2D and kinematic. A trajectory is a sequence of positions traversed at constant speed. There is no physics engine, which is a scope decision: physics would improve the pictures and change none of the numbers.

The observer is Boltzmann-rational in exactly the form of [2], equation 8, and our legibility metric is their equation 9 including the weighting $f(t) = T - t$. Both were compared line by line against the implementation. Our only addition is an inverse temperature, which is recorded in the observer's name in every result.

That posterior rests on the optimal cost-to-go, so we compute it exactly rather than approximately. Obstacles are convex polygons, the shortest obstacle-avoiding path is a polyline through their vertices, and the optimum comes from search over the visibility graph: no grid, no resolution to defend. The geometric predicate underneath is guarded, decided in floating point wherever an error bound proves the sign cannot have changed and in exact rational arithmetic otherwise. On 20,000 near-collinear cases an unguarded floating point determinant returns the wrong sign on more than a tenth of them; the guarded predicate matches rational arithmetic on all of them, and the test suite asserts both halves so the corpus cannot quietly become easy.

Keep-out zones do not block motion. A trajectory may cross one and is scored for having done so. If they blocked motion there would be no trade to measure.

Legibility optimisation is bounded by a ceiling on the cost ratio. That bound is not our idea: [2] adds a cost regulariser and [3] a hard trust region, both because a robot can otherwise make a trajectory ever more legible at ever greater cost. Sweeping the ceiling turns a single trajectory into a frontier.

3 What language models do

The prompt states the room, the goals, which goal the robot is going to, the obstacles, the keep-out zones, and a path budget in absolute units. It asks for a trajectory that makes the destination clear as early as possible, and asks the model to judge its own answer. Temperature 0.7, five samples per world, eight worlds.

Three patterns repeat across both models. Every one of the 80 decodes claimed legibility. In the world built so that the clearest signal and the keep-out zone conflict (Figure 1), both models beat the shortest path on five of five samples and both entered the zone on five of five: ten decodes out of ten bought the clarity and paid the constraint for it. And in the world with three goals where the true one lies between the others, neither model beat the shortest path on any sample, which is what [2] predicts when exaggeration towards a middle goal points at a different goal.

Over 40 decodes each	Llama 3.3 70B	Qwen 2.5 7B
replies that parsed	40	40
called legible by the model	40	40
physically possible	29	26
clearer than the shortest path	20	10
over the stated path budget	15	9
entered a keep-out zone	7	7

Table 1: Counts at a stated path budget of 25 per cent. Every decode claimed to be legible, including those that were not physically possible.




Figure 1: The world built so that the clearest signal is the one the robot may not send. The cheapest route to the true goal A (black) stays clear of the keep-out zone (hatched), so safety costs nothing until a planner tries to be legible. Every trajectory either model returned (red, ten of them, both models, five samples each) is more legible than the shortest path, and every one of the ten crosses the zone. Drawn from the committed records, not by hand.

	Llama 3.3 70B		Qwen 2.5 7B	
Stated budget	median cost	over	median cost	over
1.10	1.2817	30	1.1663	17
1.25	1.2817	15	1.1588	9
1.50	1.2817	10	1.1712	3
2.00	1.2817	0	1.1709	0

Table 2: Forty decodes at each stated budget, per model. The budget nearly doubles and the cost actually paid does not follow it.

Table 2 asks whether the budget violations mean the budget was too tight or that the budget was not attended to. It is the second. One model’s median cost moves by 0.012 across a budget that nearly

doubles and the other’s does not move at all. The falling violation counts are the line moving, not the models complying. Where the optimiser sits exactly on the budget in every row, because for a planner it is a constraint, for both models it is text.

4 Limitations

The metric is exactly specified and exactly reproducible, which is not the same as validated. Nobody has looked at these trajectories. The guidelines for evaluating social navigation ask that objective metrics be shown to measure what they purport to [4], and the only validity anchor available is the original user studies, which their authors bound to inside the trust region. Judge-free is not judge-proof. Comparing legibility frameworks against human observers remains the harder and more informative experiment [5].

Three narrower limits. With more than one goal the legibility score cannot reach 1, so a value of 0.93 is not 93 per cent of anything. Our observer has no option for a goal outside the declared set, though real observers form exactly that belief when motion becomes strange enough. And time to confidence, the one metric here with a free parameter, is stable in aggregate across admissible thresholds but not per trajectory, so no claim above rests on it.

Two models at one temperature over eight worlds is a pilot. The records are committed, one JSON object per line, and every number here is recomputed from them by a committed tool.

References

- [1] Jean-Luc Bastarache, Christopher Nielsen, and Stephen L. Smith. 2023. On Legible and Predictable Robot Navigation in Multi-Agent Environments. In *IEEE International Conference on Robotics and Automation (ICRA)*.
- [2] Anca D. Dragan, Kenton C. T. Lee, and Siddhartha S. Srinivasa. 2013. Legibility and Predictability of Robot Motion. In *Proceedings of the 8th ACM/IEEE International Conference on Human-Robot Interaction (HRI)*. 301–308.
- [3] Anca D. Dragan and Siddhartha S. Srinivasa. 2013. Generating Legible Motion. In *Robotics: Science and Systems (RSS)*.
- [4] Anthony Francis, Claudia Pérez-D’Arpino, Chengshu Li, Fei Xia, Alexandre Alahi, et al. 2025. Principles and Guidelines for Evaluating Social Robot Navigation Algorithms. *ACM Transactions on Human-Robot Interaction* (2025). Also available as arXiv preprint arXiv:2306.16740.
- [5] Sebastian Wallkötter, Mohamed Chetouani, and Ginevra Castellano. 2022. A New Approach to Evaluating Legibility: Comparing Legibility Frameworks Using Framework-Independent Robot Motion Trajectories. *arXiv preprint arXiv:2201.05765* (2022).